

# CS 4602

# Introduction to Machine Learning

Decision Trees  
(Bagging and Boosting)

Instructor: Po-Chih Kuo

# Roadmap

- Introduction and Basic Concepts
- Regression (**Error-Based Learning**)
- Bayesian Classifiers (**Probability-Based Learning**)
- Decision Trees (**Information-Based Learning**)
- KNN (**Similarity-Based Learning**)
- Linear Classifier
- Neural Networks
- Deep learning
- Convolutional Neural Networks
- RNN/Transformer
- Reinforcement Learning
- Model Selection and Evaluation
- Clustering
- Data Exploration & Dimensionality reduction

# Make a Decision

Problem: decide whether to wait for a table at a restaurant based on the following **attributes**:

1. **Alternate**: is there an alternative restaurant nearby?
2. **Bar**: is there a comfortable bar area to wait in?
3. **Fri/Sat**: is today Friday or Saturday?
4. **Hungry**: are we hungry?
5. **Patrons**: number of people in the restaurant (None, Some, Full)
6. **Price**: price range (\$, \$\$, \$\$\$)
7. **Raining**: is it raining outside?
8. **Reservation**: have we made a reservation?
9. **Type**: kind of restaurant (French, Italian, Thai, Burger)
10. **WaitEstimate**: estimated waiting time (0-10, 10-30, 30-60, >60)

# Attribute-based representations

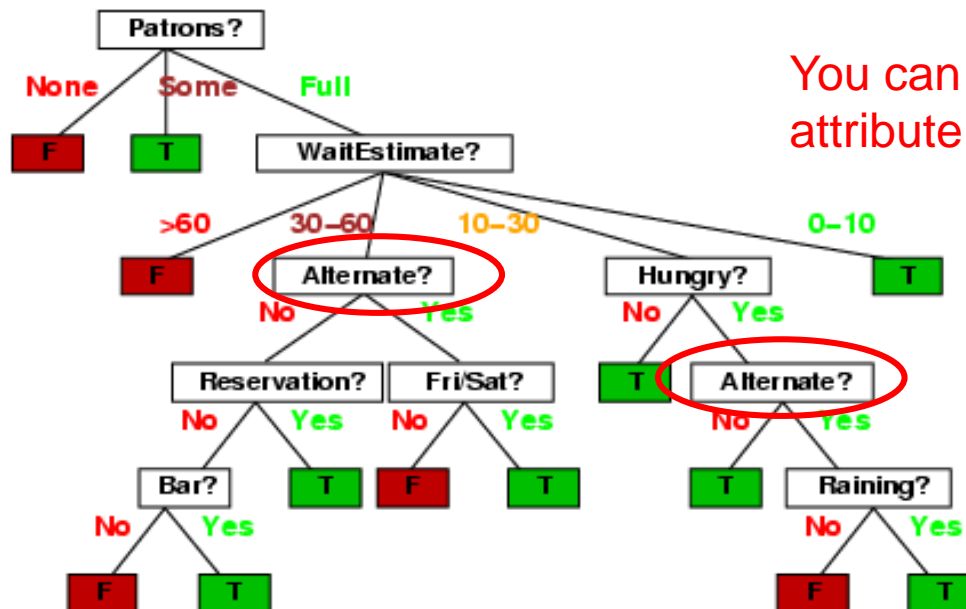
- Examples described by **attribute values** (Boolean, discrete, continuous)
- E.g., situations where I will/won't wait for a table:

Example	Attributes										Target
	Alt.	Bar	Fri	Hun	Pat	Price	Rain	Res	Type	Est.	Wait
$X_1$	T	F	F	T	Some	\$\$\$	F	T	French	0-10	T
$X_2$	T	F	F	T	Full	\$	F	F	Thai	30-60	F
$X_3$	F	T	F	F	Some	\$	F	F	Burger	0-10	T
$X_4$	T	F	T	T	Full	\$	F	F	Thai	10-30	T
$X_5$	T	F	T	F	Full	\$\$\$	F	T	French	>60	F
$X_6$	F	T	F	T	Some	\$\$	T	T	Italian	0-10	T
$X_7$	F	T	F	F	None	\$	T	F	Burger	0-10	F
$X_8$	F	F	F	T	Some	\$\$	T	T	Thai	0-10	T
$X_9$	F	T	T	F	Full	\$	T	F	Burger	>60	F
$X_{10}$	T	T	T	T	Full	\$\$\$	F	T	Italian	10-30	F
$X_{11}$	F	F	F	F	None	\$	F	F	Thai	0-10	F
$X_{12}$	T	T	T	T	Full	\$	F	F	Burger	30-60	T

- **Classification** of examples is **positive** (T) or **negative** (F)
- A number  $N$  of instances, each with attributes  $(x_1, x_2, x_3, \dots, x_d)$  and target value  $y$ .

# Decision trees


- One possible representation for hypotheses
- We imagine someone taking a sequence of decisions.
- E.g., here is the “true” tree for deciding whether to wait:



You can use the same attribute more than once.

# Expressiveness

- Decision trees can express any function of the input attributes.
- E.g., for Boolean functions, truth table row  $\rightarrow$  path to leaf:



- Trivially, there is a consistent decision tree for any training set with one path to leaf for each example (unless nondeterministic in  $x$ ) but **it probably won't generalize to new examples**
- Prefer to find more **compact** decision trees: we don't want to memorize the data, we want to find **structure** in the data!

# Hypothesis spaces

How many distinct decision trees with  $n$  Boolean attributes?


= number of Boolean functions

= number of distinct truth tables with  $2^n$  rows =  $2^{2^n}$

- E.g., with 6 Boolean attributes, how many trees?

There are 18,446,744,073,709,551,616 trees !

$X_1$	$X_2$	$Y$
0	0	?
0	1	?
1	0	?
1	1	?



$n=2$ :  $2^2 = 4$  rows. For each row we can choose T or F:  $2^4$  functions.



# Decision tree learning

- If there are so many possible trees, can we actually search this space? (solution: greedy search).
- **Aim:** find a small tree consistent with the training examples
- **Idea:** (recursively) choose "most significant" attribute as root of (sub)tree.



# Choosing an attribute

- Idea: a good attribute splits the examples into subsets that are (ideally) "all positive" or "all negative"



To wait or not to wait is still at 50%!


- Patrons or type?*

# Information theory

- **Entropy** measures the amount of uncertainty in a **probability** distribution:
  - Consider tossing a biased coin.
  - If you toss the coin VERY often, the frequency of heads is  $p$  and the frequency of tails is  $1-p$ . (fair coin  $p=0.5$ ).
  - The uncertainty in any actual outcome is given by the entropy.
  - The uncertainty is zero if  $p=0$  or  $1$  and maximal if we have  $p=0.5$ .



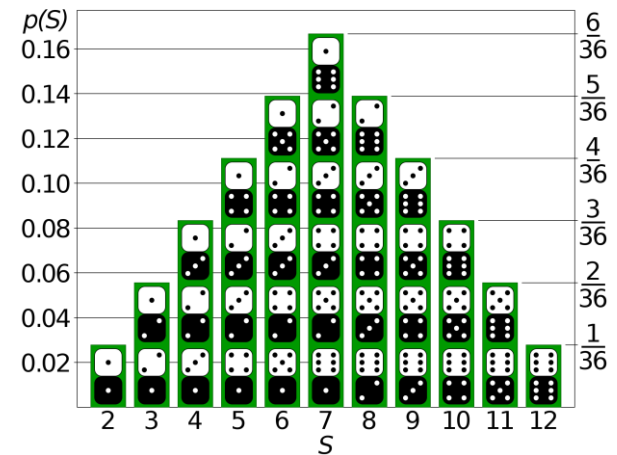
$$Entropy(X) = -p \log_2 p - (1 - p) \log_2 (1 - p)$$



# Information theory

- If there are more than two states  $s=1,2,..n$  we have (e.g. a die):

$$\begin{aligned} \text{Entropy}(X) = & \\ & -p(s=1) \log_2[p(s=1)] \\ & -p(s=2) \log_2[p(s=2)] \\ & \dots \\ & -p(s=n) \log_2[p(s=n)] \end{aligned}$$



$$\sum_{s=1}^n p(s) = 1$$

# Information theory

- Imagine we have  $p$  examples which are true (positive) and  $n$  examples which are false (negative).
- Our best estimate of true or false is given by:

$$P(\text{true}) \approx \frac{p}{p+n}$$
$$P(\text{false}) \approx \frac{n}{p+n}$$

- Hence the entropy is given by:

$$\text{Entropy}\left(\frac{p}{p+n}, \frac{n}{p+n}\right) \approx -\frac{p}{p+n} \log \frac{p}{p+n} - \frac{n}{p+n} \log \frac{n}{p+n}$$

# What about Cross entropy?

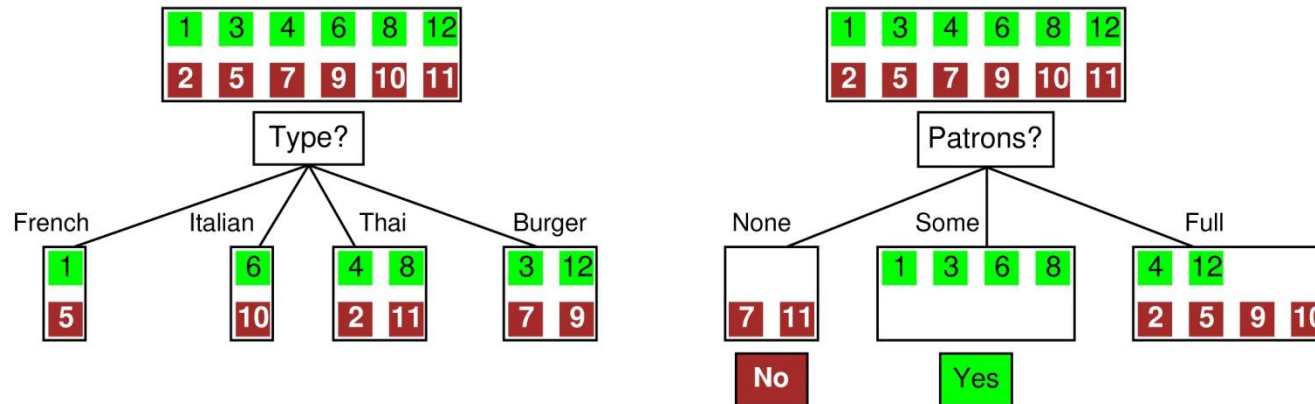
- Cross-entropy is a measure of the difference between two probability distributions.
- We'll talk about this in neural network.

# Using information Theory

- How much information do we gain if we disclose the value of some attribute?
- Answer:

uncertainty before - uncertainty after

# Example



**Before:** Entropy =  $-\frac{1}{2} \log(1/2) - \frac{1}{2} \log(1/2) = \log(2) = 1$  bit:  
There is “1 bit of information to be discovered”.

**After:** for **Type**: If we go into branch “French” we have 1 bit, similarly for the others.

French: 1bit  
Italian: 1 bit  
Thai: 1 bit  
Burger: 1bit

On average: 1 bit ! We gained nothing!

**After:** for **Patrons**: In branch “None” and “Some” entropy = 0  
In branch “Full” entropy =  $-\frac{1}{3} \log(1/3) - \frac{2}{3} \log(2/3) = 0.918....$

**So Patrons gains more information!**

# Information Gain


- How do we combine branches:

1/6 of the time we enter “None”, so we weight “None” with 1/6. Similarly: “Some” has weight: 1/3 and “Full” has weight 1/2.

$$Entropy(A) = \sum_{i=1}^n \frac{p_i + n_i}{p + n} Entropy\left(\frac{p_i}{p_i + n_i}, \frac{n_i}{p_i + n_i}\right)$$

weight for each branch

entropy for each branch.





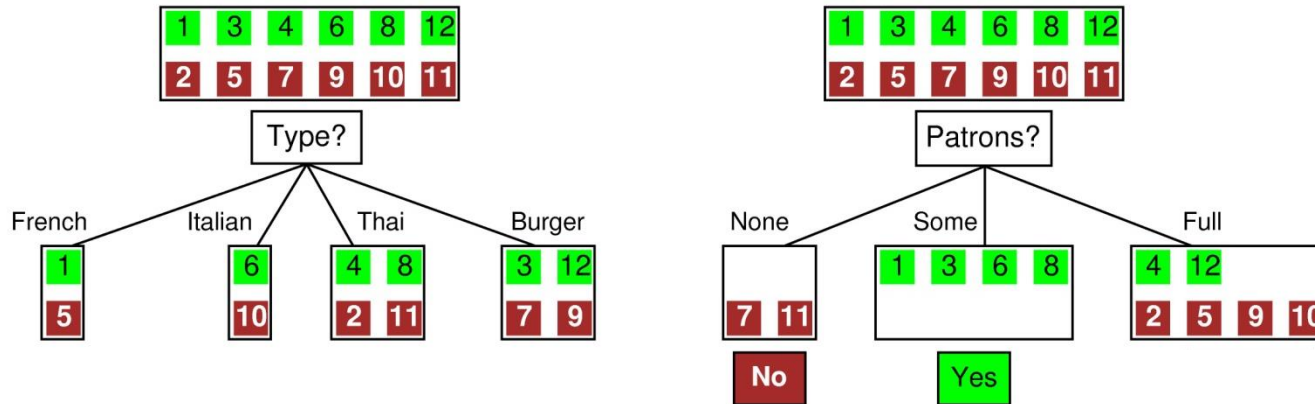
# Information Gain

- Information Gain (IG) or reduction in entropy from the attribute test:

$$IG(A) = Entropy\ before - Entropy\ after$$

- Choose the attribute with the largest IG

# Information Gain



For the training set,  $p = n = 6$ ,  $I(6/12, 6/12) = 1$  bit

$$IG(Patrons) = 1 - \left[ \frac{2}{12} I(0,1) + \frac{4}{12} I(1,0) + \frac{6}{12} I\left(\frac{2}{6}, \frac{4}{6}\right) \right] = 0.541 \text{ bits}$$

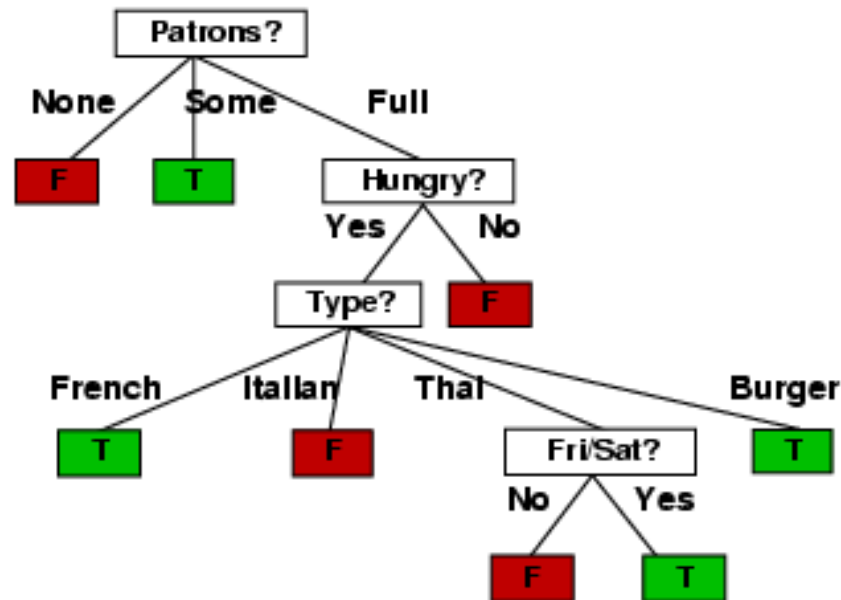
$$IG(Type) = 1 - \left[ \frac{2}{12} I\left(\frac{1}{2}, \frac{1}{2}\right) + \frac{2}{12} I\left(\frac{1}{2}, \frac{1}{2}\right) + \frac{4}{12} I\left(\frac{2}{4}, \frac{2}{4}\right) + \frac{4}{12} I\left(\frac{2}{4}, \frac{2}{4}\right) \right] = 0 \text{ bits}$$

*Patrons* has the highest IG of all attributes and so is chosen by the Decision Tree Learning algorithm as the root

Is it sensible to have the same attribute on a single branch of the tree (why)?

# Example contd.

- Decision tree learned from the 12 examples:



- Substantially simpler than “true” tree---a more complex hypothesis isn’t justified by small amount of data

# Gain-Ratio

- If 1 attribute splits in many more classes than another, it has an (unfair) advantage if we use information gain.
- The gain-ratio is designed to compensate for this problem,

$$\text{Gain - Ratio} = \frac{\text{Info - Gain}}{-\sum_{i=1}^n \frac{p_i + n_i}{p + n} \log \frac{p_i + n_i}{p + n}}$$

- if we have n uniformly populated classes the denominator is  $\log_2(n)$  which penalized relative to 1 for 2 uniformly populated classes.

# Gain-Ratio (Example)

$$\text{GainRatio}(T,X) = \frac{\text{Gain}(T,X)}{\text{SplitInformation}(T,X)}$$

$$\text{Split}(T,X) = -\sum_{c \in A} P(c) \log_2 P(c)$$

		Play Golf		
		Yes	No	total
Outlook	Sunny	3	2	5
	Overcast	4	0	4
	Rainy	2	3	5
Gain = 0.247				

$$\begin{aligned} \text{Split}(\text{Play}, \text{Outlook}) &= - (5/14 * \log_2(5/14) + 4/14 * \log_2(4/14) + 5/14 * \log_2(5/14)) \\ &= 1.577 \end{aligned}$$

$$\text{Gain Ratio}(\text{Play}, \text{Outlook}) = 0.247/1.577 = 0.156$$

# What to Do if...

- In some leaf there are no examples:

Choose True or False according to the number of positive/negative examples at your parent.

- There are no attributes left

Two or more examples have the same attributes but different label: we have an error/noise. Stop and use majority vote.

# Continuous Variables



- If variables are continuous we can bin them, or..
- We can learn a simple classifier on a single dimension
- E.g. we can find a decision point which classifies all data to the left of that point in one class and all data to the right in the other (decision stump – next slide)
- We can also use a small subset of dimensions and train a linear classifier (e.g. logistic regression classifier).

# Decision Stump

- Data:  $\{(X_1, Y_1), (X_2, Y_2), \dots, (X_n, Y_n)\}$

attributes (e.g. temperature outside) label (e.g. True or False,  
0 or 1  
**-1 or +1**)

$$h(x) = \begin{cases} +1 & \text{if } X_i > \theta \\ -1 & \text{if } X_i < \theta \end{cases}$$



So, we choose one attribute "i", a sign "+/-" and a threshold  $\theta$   
This determines a half space to be  $+1$  and the other half  $-1$ .



# A practical way for continuous variables

Ex.	Est	Wa it
X <sub>1</sub>	8	T
X <sub>2</sub>	40	F
X <sub>3</sub>	8	T
X <sub>4</sub>	12	T
X <sub>5</sub>	70	F
X <sub>6</sub>	3	T
X <sub>7</sub>	7	F
X <sub>8</sub>	6	T
X <sub>9</sub>	80	F
X <sub>10</sub>	20	F
X <sub>11</sub>	8	F
X <sub>12</sub>	40	T

sort  
→

Ex.	Est	Wa it
X <sub>6</sub>	3	T
X <sub>8</sub>	6	T
X <sub>7</sub>	7	F
X <sub>1</sub>	8	T
X <sub>3</sub>	8	T
X <sub>11</sub>	8	F
X <sub>4</sub>	12	T
X <sub>10</sub>	20	F
X <sub>2</sub>	40	F
X <sub>12</sub>	40	T
X <sub>5</sub>	70	F
X <sub>9</sub>	80	F

Middle point (what else?)


4.5 → IG  
6.5 → IG  
7.5 → IG

10 → IG  
16 → IG  
30 → IG

55 → IG  
75 → IG

if there are **N** possible values, we would have at most **N-1** possible splits.

→ Max(IGs) is the IG for this attribute (EST.)




# When to Stop ?

- If we keep going until perfect classification we might over-fit.
- Heuristics:
  - Stop when Info-Gain (Gain-Ratio) is smaller than threshold
  - Stop when there are M examples in each leaf node
- Penalize complex trees by minimizing with “complexity” = # nodes.  
Note: if tree grows, complexity grows but entropy shrinks.

$$\alpha \times \text{complexity} + \sum_{\text{all leaves}} \text{entropy}(\text{leaf})$$

- Compute many full grown trees on subsets of data and test them on hold-out data. Pick the best or average their prediction.
- Do a statistical test: is the increase in information significant.



$$\chi^2_{1,\alpha} = \frac{\left( \frac{n_1 + p_1}{n + p} n - n_1 \right)^2}{\left( \frac{n_1 + p_1}{n + p} n \right)} + \frac{\left( \frac{n_1 + p_1}{n + p} p - p_1 \right)^2}{\left( \frac{n_1 + p_1}{n + p} p \right)}$$

Degree of freedom

# How to improve Decision Trees?

**“Unity is strength!”**



**Ensemble methods!**

# Ensemble methods: bagging, boosting

## Aims

- **Bagging (bootstrap aggregating)** algorithms aim to reduce the complexity of models that **overfit** the training data.
- **Boosting** is an approach to increase the complexity of models that **underfit** the training data.

# How do they work?

- **Bagging:** learn homogeneous weak learners **independently** from each other in **parallel** and combine them following some kind of deterministic averaging process.
- **Boosting:** learn homogeneous weak learners **sequentially** in an adaptative way (a base model depends on the previous ones) and combine them following a deterministic strategy.
- **Stacking:** learn heterogeneous weak learners in **parallel** and combines them by **training a meta-model** to output a prediction based on the different weak models predictions.

Ref: <https://towardsdatascience.com/ensemble-methods-bagging-boosting-and-stacking-c9214a10a205>

# Bagging

- “bagging” (standing for “bootstrap aggregating”) that aims at producing an ensemble model that is **more robust** than the individual models composing it.
- The low correlation between models is the key.
- Bootstrapping



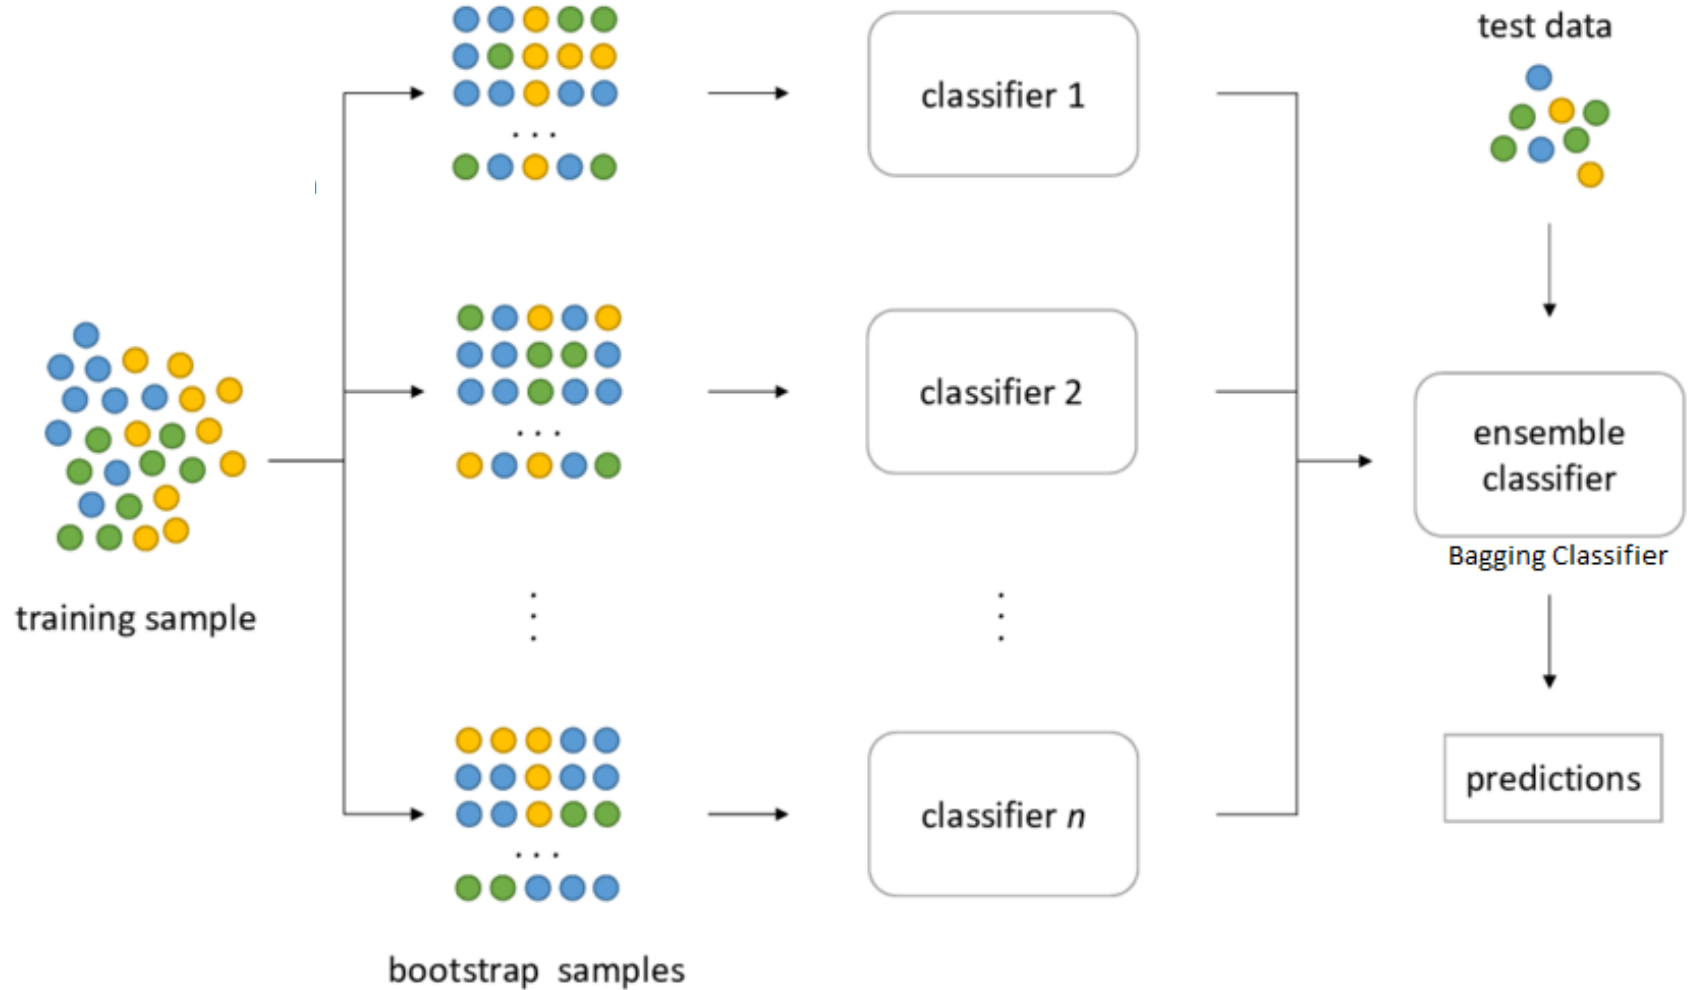
# Bagging

- Assuming that we have  $L$  bootstrap samples (approximations of  $L$  independent datasets)
- We can fit  $L$  almost independent weak learners (one on each dataset)

$$w_1(.), w_2(.), \dots, w_L(.)$$

- Then aggregate them into some kind of averaging or voting process in order to get an ensemble model with a lower variance.

$$s_L(.) = \frac{1}{L} \sum_{l=1}^L w_l(.)$$



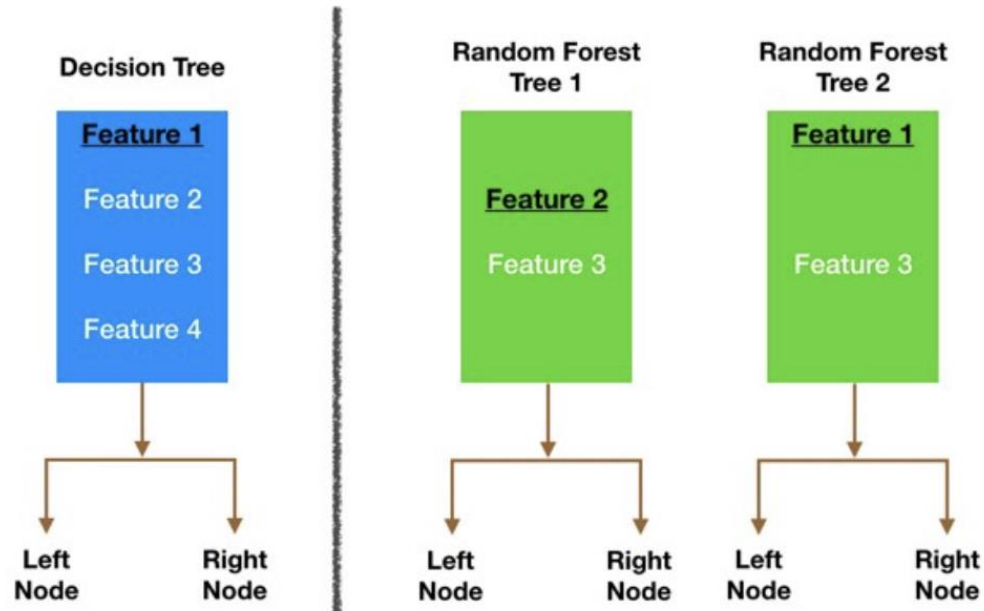
**Bagging Classifier Process Flow**



# Random forests

- When growing each tree, instead of only sampling over the observations in the dataset to generate a bootstrap sample, it also **sample over features** and keep only a random subset of them to build the tree.
  - It forces even more variation amongst the trees in the model.
  - It reduces the correlation between the different returned outputs.
  - It makes the decision making process more robust to missing data.


# Random forests



# Boosting

- Main idea:
  - Train classifiers (e.g. decision trees) in a sequence.
  - A new classifier should focus on those cases which were incorrectly classified in the last round.
  - Combine the classifiers by letting them vote on the final prediction (like bagging).
  - Each classifier could be (should be) very “weak”, e.g. a decision stump.

# Example




this line is one simple classifier saying that everything to the left “+” and everything to the right is “-”

# Boosting Intuition


- We adaptively weigh each data case.
- Data cases which are wrongly classified get high weight (the algorithm will focus on them)
- Each boosting round learns a new (simple) classifier on the weighed dataset.
- These classifiers are weighed to combine them into a single powerful classifier.
- Classifiers that obtain low training error rate have high weight.
- We stop by using monitoring a hold out set (cross-validation).

$$s_L(.) = \sum_{l=1}^L c_l \times w_l(.) \quad \longrightarrow \quad s_l(.) = s_{l-1}(.) + c_l \times w_l(.)$$

# Boosting in a Picture




# AdaBoost (Example)



Original Training set : Equal Weights to all training samples

# AdaBoost (Example)

ROUND 1



error




$$\epsilon_1 = 0.30$$

$$\alpha_1 = 0.42$$

Coef.

$D_2$




$$\alpha_t = \frac{1}{2} \ln \frac{1 - \epsilon_t}{\epsilon_t}$$




# AdaBoost (Example)

ROUND 2




$$\begin{aligned}\epsilon_2 &= 0.21 \\ \alpha_2 &= 0.65\end{aligned}$$

$D_3$



# AdaBoost (Example)

ROUND 3



# AdaBoost (Example)

$$H_{\text{final}} = \text{sign} \left( 0.42 \begin{array}{|c|} \hline \text{blue} \\ \hline \text{red} \end{array} + 0.65 \begin{array}{|c|} \hline \text{blue} \\ \hline \text{red} \end{array} + 0.92 \begin{array}{|c|} \hline \text{blue} \\ \hline \text{red} \end{array} \right)$$

The diagram illustrates the final AdaBoost hypothesis  $H_{\text{final}}$  as a weighted sum of three weak classifiers. Each classifier is represented by a square divided into blue and red regions by a decision boundary. The weights are 0.42, 0.65, and 0.92. The first two classifiers have vertical decision boundaries, and the third has a horizontal one. The entire sum is enclosed in a green sign function.

# AdaBoost (Algorithm)

Given:  $(x_1, y_1), \dots, (x_m, y_m)$  where  $x_i \in X, y_i \in Y = \{-1, +1\}$

Initialize  $D_1(i) = \frac{1}{m}$

For  $t = 1, \dots, T$  :

- Find the classifier  $h_t : X \rightarrow \{-1, +1\}$  that minimizes the error with respect to the distribution  $D_t$ :

$$h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i) [y(i) \neq h_j(x_i)]$$

- Prerequisite:  $\epsilon_t < 0.5$ , otherwise stop.
- Choose  $\alpha_t \in \mathbf{R}$  , typically  $\alpha_t = \frac{1}{2} \ln \frac{1 - \epsilon_t}{\epsilon_t}$  where  $\epsilon_t$  is the weighted error rate of classifier  $h_t$ .

- Update:  $D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$  where  $Z_t$  is a normalization factor

$$\exp(-\alpha_t y_i h_t(x_i)) \begin{cases} < 1, & y(i) = h_t(x_i) \\ > 1, & y(i) \neq h_t(x_i) \end{cases}$$

Output the final classifier:  $H(x) = \text{sign} \left( \sum_{t=1}^T \alpha_t h_t(x) \right)$

# Gradient boosting

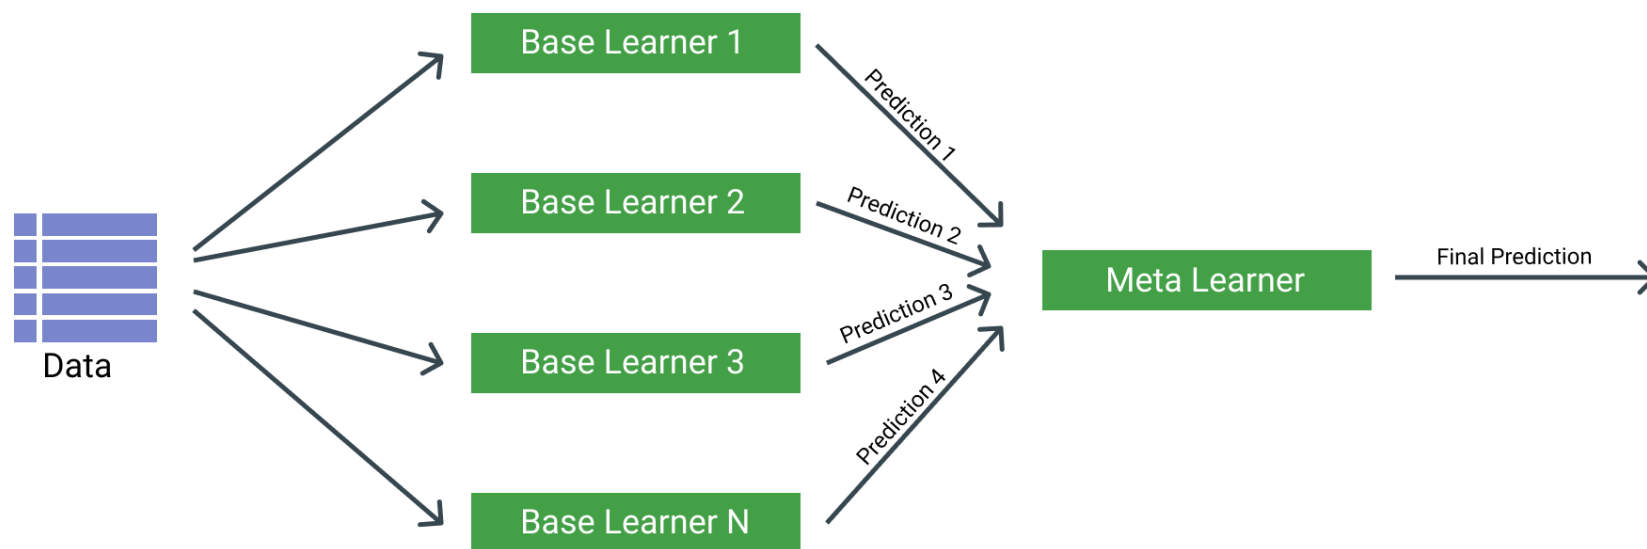
- Gradient boosting **casts the problem into a gradient descent one**: at each iteration we fit a weak learner to the opposite of the gradient of the current fitting error with respect to the current ensemble model.

$$\begin{array}{ccc} \text{Bagging} & & \text{AdaBoost} \\ s_L(.) = \sum_{l=1}^L c_l \times w_l(.) & \xrightarrow{\quad} & s_l(.) = s_{l-1}(.) + c_l \times w_l(.) \\ & & \downarrow \\ & & s_l(.) = s_{l-1}(.) - c_l \times \nabla_{s_{l-1}} E(s_{l-1})(.) \end{array}$$

# XGBoost(Extreme Gradient Boosting)

- Additional tricks that make learning much more efficient:
  - Implements **regularization** helping reduce overfit
  - Implements **parallel** processing being much faster (10x) than GB
  - Allows users to define **custom optimization objectives** and evaluation criteria
  - XGBoost has an in-built routine to handle **missing values**
  - XGBoost **prunes** the tree backwards and removes splits beyond which there is no positive gain
  - XGBoost allows a user to run a **cross-validation at each iteration** of the boosting process

# Stacking



# Questions?



model trained  
for 1000 epochs



model trained  
for 100 epochs



model trained  
for 1 epoch

<https://colab.research.google.com/github/jakevdp/PythonDataScienceHandbook/blob/master/notebooks/05.08-Random-Forests.ipynb>