

CS 4602

Introduction to Machine Learning

Neural Networks

Instructor: Po-Chih Kuo

Roadmap


- Introduction and Basic Concepts
- Regression
- Bayesian Classifiers
- Decision Trees
- KNN
- Linear Classifier
- Neural Networks
- Deep learning
- Convolutional Neural Networks
- RNN/Transformer
- Reinforcement Learning
- Model Selection and Evaluation
- Clustering
- Data Exploration & Dimensionality reduction

Outline

- Motivation
- Multilayer perceptrons (MLP)
- Backpropagation
- Extension

What are connectionist neural networks?

- **Connectionism** refers to a computer modeling approach to computation that is loosely based upon the architecture of the brain.
- Many different models, but all include:
 - Multiple, individual “**nodes**” or “**units**” that operate at the same time (in parallel)
 - A **network** that connects the nodes together
 - Information is stored in a distributed fashion among the **links** that connect the nodes
 - **Learning** can occur with gradual changes in connection strength



Neural Network History

- History traces back to the 50's
 - became popular in the 80's with work by Rumelhart, Hinton, and McClelland
 - A General Framework for Parallel Distributed Processing
- Peaked in the 90's. Today:
 - Hundreds of variants
 - Less a model of the actual brain than a useful tool, but still some debate
- Numerous applications
 - Handwriting, face, speech recognition
 - Self-driving Vehicles
 - Models of reading, sentence production, dreaming
- Debate for philosophers and cognitive scientists
 - Can human consciousness or cognitive abilities be explained by a connectionist model?

Comparison of Brains and Traditional Computers



- 200 billion neurons, 32 trillion synapses
- Element size: 10^{-6} m
- Energy use: 25W
- Processing speed: 100 Hz
- Parallel, Distributed
- Fault Tolerant
- Learns: Yes
- Conscious: Usually



- 1 billion bytes RAM but trillions of bytes on disk
- Element size: 10^{-9} m
- Energy watt: 30-90W (CPU)
- Processing speed: 10^9 Hz
- Serial, Centralized
- Generally not Fault Tolerant
- Learns: Some
- Conscious: Generally No



Biological Inspiration

Idea : To make the computer more robust, intelligent, and learn, ...
Let's model our computer software (and/or hardware) after the brain




Neurons in the Brain

- Although heterogeneous, at a low level the brain is composed of neurons
 - A neuron receives input from other neurons (generally thousands) from its synapses
 - Inputs are approximately summed
 - When the input exceeds a threshold the neuron sends an electrical spike that travels that travels from the body, down the axon, to the next neuron(s)



Perceptrons

- Initial proposal of connectionist networks
- Rosenblatt, 50's and 60's
- Essentially a linear discriminant composed of nodes, weights




Activation Function

$$O = \begin{cases} 1 : \left(\sum_i w_i x_i \right) + w_0 > 0 \\ 0 : \text{otherwise} \end{cases}$$


Linear separability

- Consider a single-layer perceptron
 - Assume threshold units
 - Assume binary inputs and outputs
 - Weighted sum forms a linear hyperplane $\sum_i w_i x_i = 0$
- Consider a single output network with two inputs
 - Only functions that are linearly separable can be computed
 - Example: AND is linearly separable




Linear inseparability

- Single-layer perceptron with threshold units fails if problem is not linearly separable
 - Example: XOR





Solution in 1980s: Multilayer perceptrons

- Removes many limitations of single-layer networks
 - Can solve XOR
- How to Draw a two-layer perceptron that computes the XOR function?
 - 2 binary inputs x_1 and x_2
 - 1 binary output
 - One “hidden” layer
 - Find the appropriate weights and threshold





Multilayer perceptrons

- Examples of two-layer perceptrons that compute XOR




- E.g. Right side network
 - Output is 1
if and only if $x_1 + x_2 - 2(x_1 + x_2 - 1.5 > 0) - 0.5 > 0$

Multilayer perceptron




The most common output function (Sigmoid):

$$\Psi(a) = \frac{1}{1 + e^{-\beta a}}$$





(non-linear function)





Source: <http://colinfahey.com/>



Perceptrons as Constraint Satisfaction Networks





Perceptrons as Constraint Satisfaction Networks





Perceptrons as Constraint Satisfaction Networks



Perceptrons as Constraint Satisfaction Networks







Perceptrons as Constraint Satisfaction Networks







Decision Boundary

- 0 hidden layers: linear classifier
 - Hyperplanes




Decision Boundary




- 1 hidden layer
 - Boundary of convex region (open or closed)



Decision Boundary



- 2 hidden layers
 - Combinations of convex regions



Outline

- Motivation
- Multilayer perceptrons (MLP)
- **Backpropagation (BP)**
- Extension

What about learning?


- Perceptron learning rule

$$\mathbf{w}' = \mathbf{w} + \alpha \sum_{n \in M} \mathbf{x}^n (Y_d - Y)$$

Table 6.3 Example of perceptron learning: the logical operation AND


Epoch	Inputs		Desired output	Initial weights		Actual output	Error	Final weights	
	x_1	x_2	Y_d	w_1	w_2	Y	e	w_1	w_2
1	0	0	0	0.3	-0.1	0	0	0.3	-0.1
	0	1	0	0.3	-0.1	0	0	0.3	-0.1
	1	0	0	0.3	-0.1	1	-1	0.2	-0.1
	1	1	1	0.2	-0.1	0	1	0.3	0.0
2	0	0	0	0.3	0.0	0	0	0.3	0.0
	0	1	0	0.3	0.0	0	0	0.3	0.0
	1	0	0	0.3	0.0	1	-1	0.2	0.0
	1	1	1	0.2	0.0	1	0	0.2	0.0
3	0	0	0	0.2	0.0	0	0	0.2	0.0
	0	1	0	0.2	0.0	0	0	0.2	0.0
	1	0	0	0.2	0.0	1	-1	0.1	0.0
	1	1	1	0.1	0.0	0	1	0.2	0.1
4	0	0	0	0.2	0.1	0	0	0.2	0.1
	0	1	0	0.2	0.1	0	0	0.2	0.1
	1	0	0	0.2	0.1	1	-1	0.1	0.1
	1	1	1	0.1	0.1	1	0	0.1	0.1
5	0	0	0	0.1	0.1	0	0	0.1	0.1
	0	1	0	0.1	0.1	0	0	0.1	0.1
	1	0	0	0.1	0.1	0	0	0.1	0.1
	1	1	1	0.1	0.1	1	0	0.1	0.1

Threshold $\theta = 0.2$, learning rate $\alpha = 0.1$




Learning networks

- We want networks that configure themselves
 - Learn from the input data or from training examples
 - Generalize from learned data



Can this network configure itself to solve a problem?

How do we train it?



Gradient-descent learning: Use a **differentiable** activation function
(Refer to previous lecture slides)

Perceptron vs. Gradient Descent Rule

Perceptron learning rule guaranteed to succeed if


- Training examples are linearly separable
- Sufficiently small learning rate η

Linear unit training rules uses gradient descent


- Use a differentiable activation function
- Guaranteed to converge to hypothesis with minimum squared error
- Given sufficiently small learning rate η
- Even when training data contains noise
- Even when training data not separable by H

Neuron with Sigmoid Function

inputs



Gradient Descent Rule for Sigmoid Output Function



$$E^p[w_1, \dots, w_n] = \frac{1}{2} (t^p - y^p)^2$$

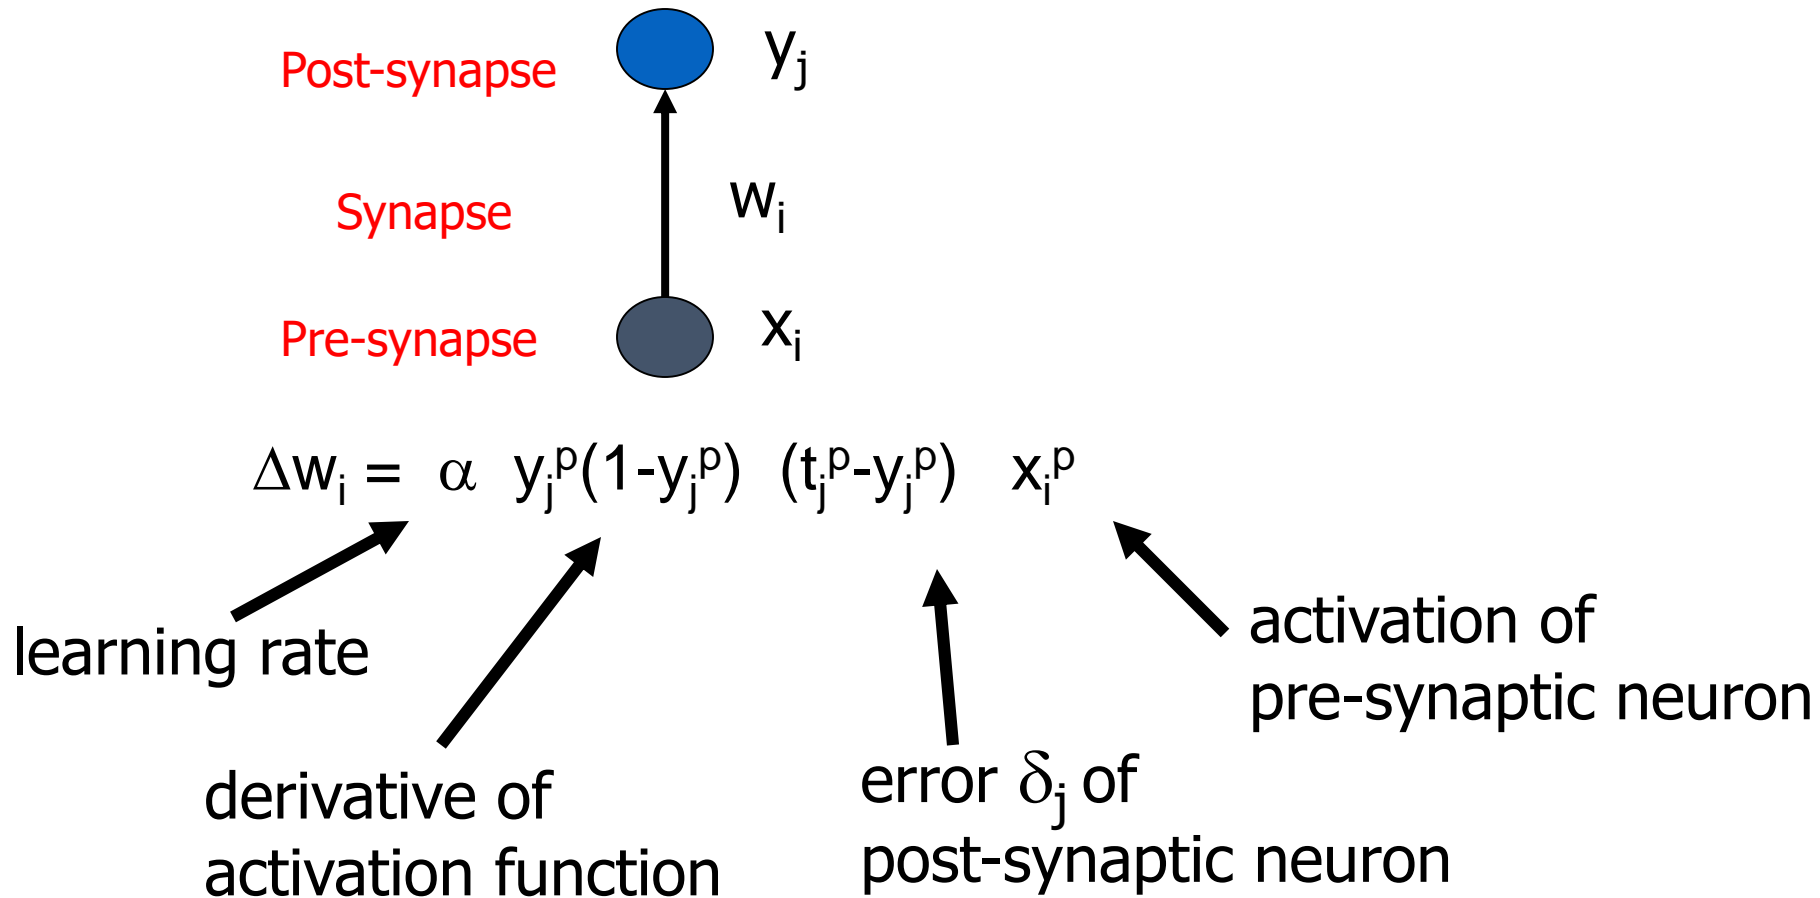
$$\begin{aligned} \frac{\partial E^p}{\partial w_i} &= \frac{\partial}{\partial w_i} \frac{1}{2} (t^p - y^p)^2 \\ &= \frac{\partial}{\partial w_i} \frac{1}{2} (t^p - \sigma(\sum_i w_i x_i^p))^2 \\ &= (t^p - y^p) \sigma'(\sum_i w_i x_i^p) (-x_i^p) \end{aligned}$$

$$\text{for } y = \sigma(a) = 1/(1+e^{-a})$$

$$\sigma'(a) = e^{-a}/(1+e^{-a})^2 = \sigma(a) (1 - \sigma(a))$$

$$\Delta w_i = \alpha y^p (1 - y^p) (t^p - y^p) x_i^p$$

Gradient Descent Learning Rule




Learning with hidden units

- Networks without hidden units are very limited in the input-output mappings they can model.
- More layers of linear units do not help. It's still linear.
- We need multiple layers of adaptive non-linear hidden units. This gives us a universal approximator. But how can we train such nets?
 - We need an efficient way of adapting **all** the weights, not just the last layer.
 - Learning the weights going into hidden units is equivalent to learning **features**.
 - It's hard to tell directly what hidden units should do.

Learning by perturbing weights

- Randomly perturb **one** weight and see if it improves performance. If so, save the change.
 - **Very inefficient.** We need to do multiple forward passes to change one weight.
- Randomly perturb **all** the weights in parallel and correlate the performance gain with the weight changes.
 - We need lots of trials to “see” the effect of changing one weight through the noise created by all the others.

Learning the hidden to output weights is **easy**.




Learning the input to hidden weights is **hard**.

The idea behind backpropagation

- We don't know what the hidden units ought to do, but we can compute how fast the error changes as we change a hidden activity.
 - Instead of using desired activities to train the hidden units, use **error derivatives w.r.t. hidden activities**.
 - Each hidden activity can affect many output units and can therefore have many separate effects on the error. These effects must be combined.
 - We can compute error derivatives for **all** the hidden units efficiently.
 - Once we have the error derivatives for the hidden activities, it's easy to get the error derivatives for the weights going into a hidden unit.

Multi-Layer Networks



Training Rule for Weights to the Output Layer

$$E^p[w_{ij}] = 1/2 \sum_j (t_j^p - y_j^p)^2$$

$$\begin{aligned} \partial E^p / \partial w_{ij} &= \partial / \partial w_{ij} 1/2 \sum_j (t_j^p - y_j^p)^2 \\ &= \dots \\ &= - y_j^p (1 - y_j^p) (t_j^p - y_j^p) x_i^p \end{aligned}$$

Derivative of activation function


Learning rate

Error of post-synaptic neuron


$$\begin{aligned} \Delta w_{ij} &= \alpha y_j^p (1 - y_j^p) (t_j^p - y_j^p) x_i^p \\ &= \alpha \delta_j^p x_i^p \end{aligned}$$

Activation of pre-synaptic neuron

with $\delta_j^p := y_j^p (1 - y_j^p) (t_j^p - y_j^p)$



Training Rule for Weights to the Hidden Layer




$$E^p[w_{ik}] = \frac{1}{2} \sum_j (t_j^p - y_j^p)^2$$

$$\begin{aligned} \frac{\partial E^p}{\partial w_{ik}} &= \frac{\partial}{\partial w_{ik}} \frac{1}{2} \sum_j (t_j^p - y_j^p)^2 \\ &= \frac{\partial}{\partial w_{ik}} \frac{1}{2} \sum_j (t_j^p - \sigma(\sum_k w_{kj} x_k^p))^2 \\ &= \frac{\partial}{\partial w_{ik}} \frac{1}{2} \sum_j (t_j^p - \sigma(\sum_k w_{kj} \sigma(\sum_i w_{ik} x_i^p)))^2 \\ &= -\sum_j (t_j^p - y_j^p) \sigma'_j(a) w_{jk} \sigma'_k(a) x_i^p \\ &= -\sum_j \delta_j w_{kj} \sigma'_k(a) x_i^p \\ &= -\sum_j \delta_j w_{kj} x_k (1 - x_k) x_i^p \end{aligned}$$

$$\Delta w_{ik} = \alpha \delta_k x_i^p$$

$$\text{with } \delta_k = \sum_j \delta_j w_{kj} x_k (1 - x_k)$$

Backpropagation




Backward step:
propagate errors from
output to hidden layer

Forward step:
Propagate activation
from input to output layer

Backpropagation Algorithm

- Initialize each w_i to some small random value
- Until the termination condition is met, Do
 - For each training example $\langle (x_1, \dots, x_n), t \rangle$ Do
 - Input the instance (x_1, \dots, x_n) to the network and compute the network outputs y_k **Forward Pass**
 - For each output unit k
$$\delta_k = y_k(1 - y_k)(t_k - y_k)$$
 Backward Pass
 - For each hidden unit h
$$\delta_h = y_h(1 - y_h) \sum_k w_{h,k} \delta_k$$
 - For each network weight $w_{i,j}$ Do **Update**
$$w_{i,j} = w_{i,j} + \Delta w_{i,j} \quad \text{where } \Delta w_{i,j} = \eta \delta_j x_{i,j}$$

BP: A worked example




Current state:

- Weights on arrows e.g. $w_{13} = 3$, $w_{35} = 2$, $w_{24} = 5$
- Bias weights, e.g. bias for unit 4 (u_4) is $w_{04} = -6$

Training example (e.g. for logical OR problem):

- Input pattern is $x_1 = 1$, $x_2 = 0$
- Target output is $y_{\text{target}} = 1$

Worked example: Forward Pass



Output for any neuron/unit j can be calculated from:

$$a_j = \sum_i w_{ij} x_i$$


$$y_j = f(a_j) = \frac{1}{1 + e^{-a_j}}$$

e.g Calculating output for Neuron/unit 3 in hidden layer:


$$a_3 = 1*1 + 3*1 + 4*0 = 4$$

$$y_3 = f(4) = \frac{1}{1 + e^{-4}} = 0.982$$

Worked example: Forward Pass



Worked example: Backward Pass



Now compute delta values starting at the output:


$$\begin{aligned}\delta_5 &= y_5(1 - y_5)(y_{\text{target}} - y_5) \\ &= 0.51(1 - 0.51) \times 0.49 \\ &= \mathbf{0.1225}\end{aligned}$$

Then for hidden units:

$$\begin{aligned}\delta_4 &= y_4(1 - y_4) w_{45} \delta_5 \\ &= 0.5(1 - 0.5) \times 4 \times 0.1225 \\ &= \mathbf{0.1225}\end{aligned}$$

$$\begin{aligned}\delta_3 &= y_3(1 - y_3) w_{35} \delta_5 \\ &= 0.982(1 - 0.982) \times 2 \times 0.1225 \\ &= \mathbf{0.0043}\end{aligned}$$

Worked example: Update Weights



- ◆ Set learning rate $\eta = 0.1$
Change weights by:

$$\Delta w_{ij} = \eta \delta_j y_i$$

- ◆ e.g. bias weight on u_3 :

$$\begin{aligned} \Delta w_{03} &= \eta \delta_3 x_0 \\ &= 0.1 * 0.0043 * 1 \\ &= 0.0004 \end{aligned}$$

$$\begin{aligned} \text{So, new } w_{03}' &= \\ w_{03}(\text{old}) + \Delta w_{03} \\ &= 1 + 0.0004 = 1.0004 \end{aligned}$$


- ◆ and likewise:

$$\begin{aligned} w_{13}' &= 3 + 0.0004 \\ &= 3.0004 \end{aligned}$$

For the all weights w_{ij} :

i	j	w_{ij}	δ_j	y_i	Updated w_{ij}
0	3	1	0.0043	1.0	1.0004
1	3	3	0.0043	1.0	3.0004
2	3	4	0.0043	0.0	4.0000
0	4	-6	0.1225	1.0	-5.9878
1	4	6	0.1225	1.0	6.0123
2	4	5	0.1225	0.0	5.0000
0	5	-3.92	0.1225	1.0	-3.9078
3	5	2	0.1225	0.9820	2.0120
4	5	4	0.1225	0.5	4.0061

Verification that it works



On next forward pass:

The new activations are:

$$y_3 = f(4.0008) = 0.9820$$

$$y_4 = f(0.0245) = 0.5061$$

$$y_5 = f(0.0955) = \mathbf{0.5239}$$

Thus the new error

$$(y_{\text{target}} - y_5) = (1 - 0.5239) = 0.476$$

has been reduced by **0.014**
(from **0.490** to **0.476**)

Ref: "Neural Network Learning & Expert Systems" by Stephen Gallant

Backpropagation

- Gradient descent over entire *network* weight vector
- Easily generalized to arbitrary directed graphs
- Will find a local, not necessarily global error minimum
- Often include weight *momentum* term

$$\Delta w_{i,j}(n) = \eta \delta_j x_{i,j} + \alpha \Delta w_{i,j}(n-1)$$


- Minimizes error training examples
 - Will it generalize well to unseen instances (over-fitting)?
- Training can be slow typical 1000-10000 iterations

Outline

- Motivation
- Multilayer perceptrons (MLP)
- Backpropagation
- **Extension**

Radial Basis Function Networks

output neurons




*one layer of
hidden neurons*

input nodes

Radial Basis Function Networks

output neurons




propagation function:

$$a_j = \sqrt{\sum_{i=1}^n (x_i - \mu_{i,j})^2}$$


input nodes

Radial Basis Function Networks

output neurons



*output function:
(Gauss' bell-shaped function)*




$$h(a) = e^{-\frac{a^2}{2\sigma^2}}$$

input nodes

Radial Basis Function Networks

output neurons



output of network:


$$\text{out}_j = \sum_i w_{i,j} h_i$$

input nodes

RBF networks

- Radial basis functions
 - Hidden units store means and variances
 - Hidden units compute a Gaussian function of inputs x_1, \dots, x_n that constitute the input vector \mathbf{x}
- Learn weights w_i , means μ_i , and variances σ_i by minimizing squared error function (gradient descent learning)


Radial Basis Function Network




$$h_i = \exp\left[-\frac{(\mathbf{x} - \mathbf{u}_i)^T (\mathbf{x} - \mathbf{u}_i)}{2\sigma_i^2}\right], \quad y = \sum_i h_i w_i$$

RBF Networks and Multilayer Perceptrons


output neurons



MLP:




RBF:



input nodes

Recurrent networks




- Employ feedback (positive, negative, or both)
 - Not necessarily stable
 - Symmetric connections can ensure stability
- Why use recurrent networks?
 - Can learn temporal patterns (time series or oscillations)
 - Biologically realistic
 - Majority of connections to neurons in cerebral cortex are feedback connections from local or distant neurons
- Examples
 - Hopfield network
 - Boltzmann machine (Hopfield-like net with input & output units)
 - Recurrent backpropagation networks: for small sequences, unfold network in time dimension and use backpropagation learning

Recurrent networks (con't)

- Example
 - Elman networks
 - Partially recurrent
 - Context units keep internal memory of part inputs
 - Fixed context weights
 - Backpropagation for learning
 - E.g. Can disambiguate $A \rightarrow B \rightarrow C$ and $C \rightarrow B \rightarrow A$

Elman network (1990)



Summary: Biology and Neural Networks

- So many similarities
 - Information is contained in synaptic connections
 - Network learns to perform specific functions
 - Network generalizes to new inputs
- But NNs are woefully inadequate compared with biology
 - Simplistic model of neuron and synapse, implausible learning rules
 - Hard to train large networks
 - Network construction (structure, learning rate etc.) is a heuristic art
- One obvious difference: Spike representation
 - Recent models explore spikes and spike-timing dependent plasticity
- Other Recent Trends: Probabilistic approach
 - NNs as Bayesian networks (allows principled derivation of dynamics, learning rules, and even structure of network)
 - Not clear how neurons encode probabilities in spikes

Questions?

"I work with **models**."

Others:



Me:

